

RRTO: A High Performance Transparent Offloading System for Model Inference on Robotic IoT

Author Names Omitted for Anonymous Review. Paper-ID [23]

Abstract—Fundamental robotic tasks, such as object identification and robot control, increasingly rely on Machine Learning (ML) models deployed on wireless robots, on which the heavy computation, resulting from the large number of parameters and complex operations (e.g., matrix multiplication, convolution, pooling) in ML models, is often offloaded to powerful GPU devices (e.g., cloud data centers) via the Internet of Things for these robots (robotic IoT) to achieve fast and energy-efficient inferences. Existing computation offloading systems can be categorized into two types: transparent and untransparent methods, depending on whether source code modification is required to enable offloading schedules. While untransparent offloading systems have been widely successful across various ML models, they also pose significant obstacles for deployment on robots, requiring considerable coding effort to modify the source code for each application and is inapplicable to closed-source applications. Meanwhile, existing transparent offloading methods can offload the function call of each operator within the ML models via Remote Procedure Call (RPC) to avoid source code modifications, but this one-by-one handling of RPCs leads to substantial communication costs in robotic IoT.

We present RRTO, the first high performance transparent offloading system optimized for ML model inference on robotic IoT with a novel record/replay mechanism. Recognizing that the operators invoked by ML models typically follow a fixed order, RRTO automatically records the order of operators invoked by the ML model and replays the execution of these fixed-order operators during model inference. In this way, RRTO calls all operators for once inference in once RPC, circumventing the inherent communication costs caused by existing transparent offloading mechanism and applying successful scheduling algorithms from untransparent offloading to transparent offloading. Evaluations demonstrate that RRTO improved the performance of our real-world robotic application by 4.9x to 48.5x and saved 8% to 53% power on our robot compared to other baselines without modifying any source code, achieving similar high performance as state-of-the-art untransparent offloading methods.

I. INTRODUCTION

The rapid advancement of machine learning (ML) methods has achieved remarkable success in various fundamental robotic tasks, such as object detection[18, 24, 25], robot control[20, 36, 41], and environmental perception[3, 21, 39]. Deploying these ML methods onto real-world robots typically necessitates additional hardware support due to the computationally-intensive nature of ML models (e.g., the large number of parameters, complex operations). However, directly integrating computing accelerators (e.g., GPU[31], FPGA[30], SoC[17]) onto real-world robots not only introduces additional economic costs, but also leads to increased energy consumption (e.g., 62% for KAPAO[25] on our robot), shown as "Local" in Fig 2. Consequently, instead of performing

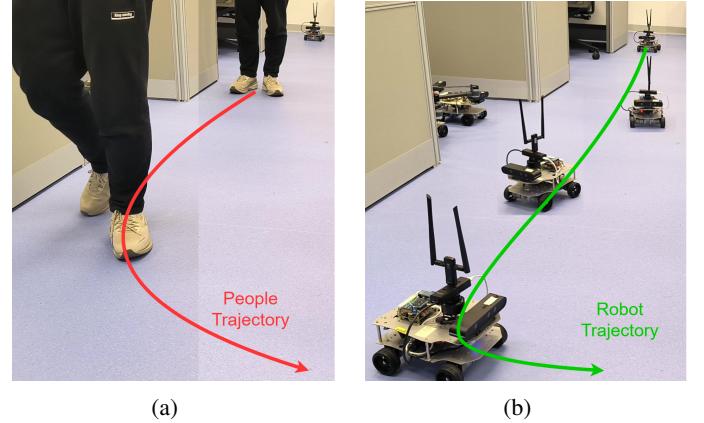


Fig. 1: A real-time people-tracking robotic application on our robot based on a well-known human pose estimation ML model, KAPAO[25].

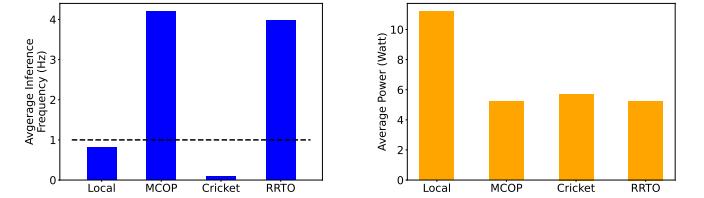


Fig. 2: Comparison between RRTO and the baselines on the application in Fig 1. The higher frequency indicates superior performance. "Local" refers to inference executed using an integrated computing accelerator on the robot, while "MCOP"[38] is a state-of-the-art (SOTA) untransparent offloading system, and "Cricket"[9] is a SOTA transparent offloading system.

computations locally, these ML applications often offload the computation of ML models to cloud data centers or edge devices equipped with powerful GPUs through the Internet of Things for these robots (robotic IoT).

Existing computation offloading systems can be categorized into two types: transparent and untransparent methods, depending on whether source code modification is required to enable offloading schedules. **Untransparent** offloading systems, such as MCOP[38], have demonstrated success across various ML models. MCOP adaptively schedules the computation of ML models to be offloaded to cloud servers based on the model's workload, network bandwidth, and the computing

power of cloud servers. This approach successfully improved the performance of our real-world robotic application by 5.2x and saved 54% power on our robot in our experiments, shown as "MCOP" in Fig 2. Such untransparency provides them with a simple and efficient scheduling method but demands significant coding effort to modify the source code for each application and can not be used on closed-source applications.

Transparent offloading methods, such as Cricket[9], provide a more convenient, but less performance approach to offload computation to the cloud servers. ML model inference consists of a series of operators (e.g., addition, convolution). Transparent offloading methods intercepts each operator's call to the corresponding system functions (e.g., `torch.add()`, `torch.convolution()`) for PyTorch[32]) and offloads all these calls to the cloud servers through Remote Procedure Call (RPC). In this way, they avoids modifications to the source code by intercepting at the system layer but has to offload the RPC calls to the cloud servers one by one and add one Round-Trip Time (RTT) of RPC to the completion time of each operator, bringing extra communication cost.

Moreover, such inherent communication cost caused by the transparent offloading mechanism becomes substantial in robotic IoT networks, rendering those outstanding scheduling algorithms ineffective. ML models commonly have hundreds of operators (e.g., 418 for KAPAO[25]), and each operator usually requires several RPC functions to complete based on different ML frameworks (e.g., `cudaGetDevice`, `cudaLaunchKernel` for PyTorch[25]), resulting in hundreds or thousands of RPC calls for a single inference (e.g., 4756 in our experiments). In robotic IoT networks, which typically rely on wireless communication for robots, each RPC RTT usually takes a few milliseconds[34] (average 2.6 milliseconds in our experiments). Consequently, transparent offloading in robotic IoT networks, shown as "Cricket" in Fig 2, leads to prolonged communication time (98% inference time in our experiments) and may even increase the final energy consumption per inference, despite reducing power on the robot by offloading.

Our key observation is that existing transparent offloading methods suffer from significant communication costs because they only invoke RPCs when an operator is utilized during the inference process, rather than anticipating their use in advance. This is because these methods are designed for general applications to leverage remote GPUs, rather than specifically for ML model inference. Such communication costs are unavoidable in general applications, as they cannot predict future operators from the upper-layer applications and thus cannot call RPCs preemptively. Fortunately, in the context of ML model inference, we find that the operators invoked by ML models typically follow a fixed order and can be predicted in advance, allowing for a more efficient inference process.

Based on this observation, we present **RRTO**, a **Transparent Offloading** system for model inference on robotic IoT with a novel **Record/Replay** mechanism: record the order of operators invoked by ML model automatically and replay the execution of these fixed-order operators during model inference. In this way, RRTO calls all operators for once inference in once

RPC, without waiting for the operator to be called via RPCs during inference. Based on this mechanism, RRTO dramatically reduces the inherent communication cost caused by traditional transparent offloading mechanism and achieves nearly the same communication cost as untransparent offloading methods, allowing the outstanding scheduling algorithms in untransparent offloading methods to be utilized in transparent offloading.

However, it is not trivial to identify the specific operators invoked during each inference, as RRTO will not be able to obtain the correct inference result even with a discrepancy of just one operator more or less. As a transparent offloading system, RRTO can only know the function calls within the upper-layer applications and must find the correct operator sequence throughout the log records without the help of hints from these applications, presenting a specific challenge for RRTO in addressing this specific issue.

To address this challenge, RRTO proposes a novel algorithm called *Data Dependency Search* to identify which operators are invoked for each inference. This algorithm first constructs a relationship graph based on the data dependencies between operators (i.e., the calculation result of the previous operator serves as the input for the next operator). It then takes operators that do not depend on any other operators as starting points and operators that do not have any dependencies on them as ending points. The algorithm subsequently searches for the longest covering operator sequence between starting and ending operators and verifies whether such an operator sequence can constitute a complete model inference process (i.e., the entire log records of operators can be covered by repeating this sequence). A specific observation that the final computation result needs to be sent back from GPU to CPU aids this algorithm in identifying the correct ending operator and dramatically reduces its search space.

We implemented RRTO on the Cricket's codebase[9] and also incorporated the outstanding scheduling algorithms of MCOP[38] into RRTO. We evaluated RRTO on a real-world Jetson Xavier NX[29] robot capable of GPU accelerated computation with a robotic application that tracks people in real time[25]. We compared RRTO with local computation, a SOTA untransparent offloading method (MCOP[38]) and a SOTA transparent offloading method (Cricket[9]) when offload computation to different GPU devices (namely edge devices with high bandwidth and cloud servers with limited bandwidth). Evaluation shows that:

- RRTO is fast. It reduces inference time by 80% to 98% compared to other baselines, similar to the reduction achieved by the SOTA untransparent offloading method.
- RRTO is energy-efficient. It saves 91% to 98% in energy consumption per inference compared to other baselines, similar to the savings of the SOTA untransparent offloading method.
- RRTO is robust in various robotic IoT environments. When the robotic IoT environment (the bandwidth to the GPU devices and the computing power of the GPU devices) changes, RRTO's superior performance remains

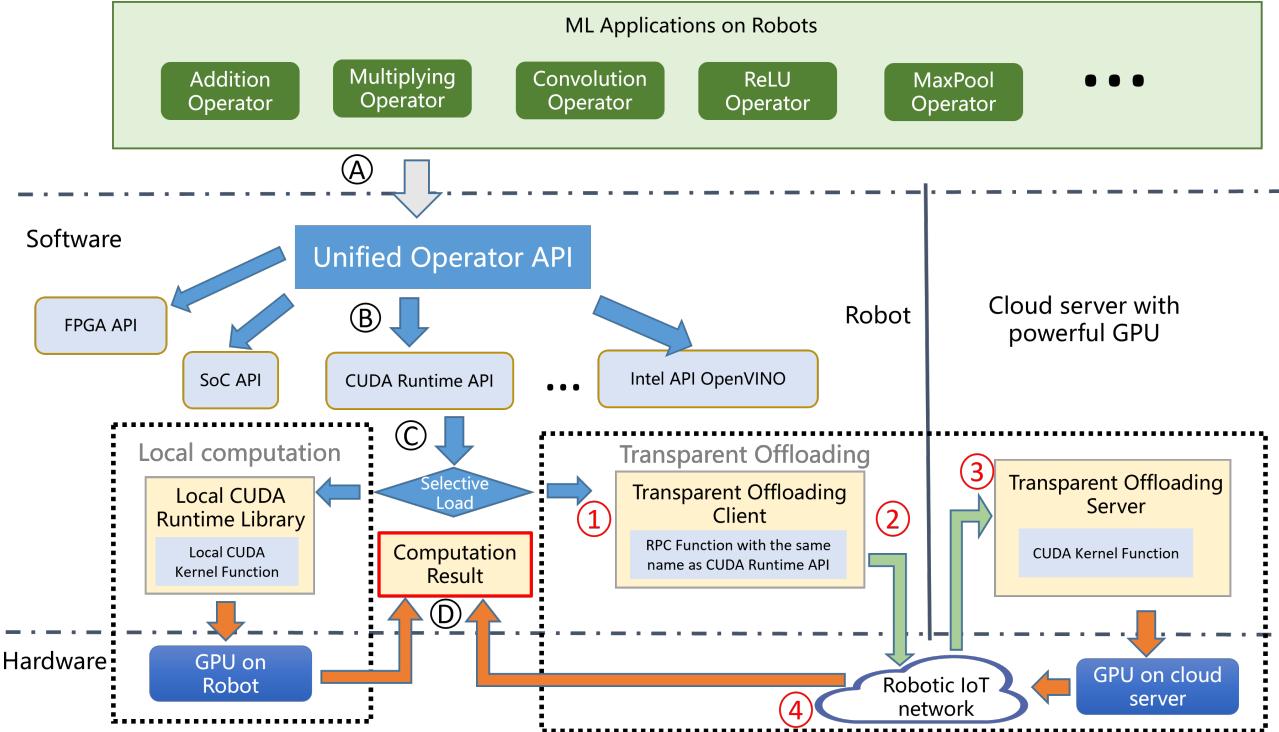


Fig. 3: Workflow of Transparent Offloading System for Model Inference on Robotic IoT

consistent as the SOTA untransparent offloading method.

Our main contribution is RRTTO, the first high-performance transparent offloading system designed for model inference on robotic IoT. RRTTO dramatically reduces the inherent communication cost caused by the traditional transparent offloading mechanism via the novel record/replay mechanism, achieving the same high performance as the SOTA untransparent offloading method without modifying any source code. We envision that RRTTO will foster the deployment of diverse robotic tasks on real-world robots in the field by providing fast, energy-efficient, and easy-to-use inference capabilities. RRTTO’s code is released on <https://github.com/rss24p23/RRTTO>

In the rest of this paper, we introduce the background of this paper in Sec. II, give an overview of RRTTO in Sec. III, present the detailed design of RRTTO in Sec. IV, evaluate RRTTO in Sec. VI, and finally conclude in Sec. VII.

II. BACKGROUND

A. Workflow of Transparent offloading

When a robot performs GPU computations locally, the system call flow diagram of the entire application is shown in the left part in Figure 3:

- A The robot application completes the entire computation process for each service by sequentially calling different operators.
- B Based on the application’s running device (GPU), each operator passes through a unified operator API to find

the local CUDA runtime library (NVIDIA GPUs provide high-performance parallel computing capabilities to applications using the CUDA runtime library[27]).

- C The local CUDA runtime API is loaded by default.
- D The robot’s local CUDA library launches the corresponding CUDA kernel functions on the robot’s GPU and returns the computation results to the upper-layer application.

Transparent offloading methods[9, 14, 23] usually takes the approach of rewriting dynamic link libraries, defining functions with the same name and using the `LD_PRELOAD` environment variable to prioritize loading the custom dynamic link library. The dynamic linker will then parse the original library function as the custom library function, thereby achieving library function interception. Subsequently, by modifying the management of GPU memory and the launch of CUDA kernel functions, the computation-related data and specific parameters of the corresponding kernel functions are sent to the remote server via RPC, realizing GPU computing transparent offloading. The primary modification occurs in step C. Similar to completing computations locally using the robot’s GPU, after steps A and B, each operator’s call to the corresponding kernel functions is intercepted by the functions with same name in the dynamic link library and offloaded to the cloud server to execute. The detailed steps are shown in the right part in Figure 3:

- (1) By modifying the dynamic link library, each operator

prioritizes calling the RPC functions with the same name as the CUDA runtime API in transparent offloading client, thereby identifying and intercepting all CUDA kernel function calls.

- (2) The transparent offloading client transmits the called CUDA runtime API and required parameters to the cloud server through the robotic IoT network via RPC.
- (3) The transparent offloading server launches the corresponding CUDA kernel functions on the cloud GPU and completes the respective computation.
- (4) The transparent offloading server sends the computation results back to the client and returns the results to the upper-layer application.

B. Related Work

Model Compression. Quantization and model distillation are the two most commonly used methods of ML model compression on the robots. Quantization[7, 13, 15] is a technique that reduces the numerical precision of model weights and activations, thereby minimizing the memory footprint and computational requirements of deep learning models. This process typically involves converting high-precision (e.g., 32-bit) floating-point values to lower-precision (e.g., 8-bit) fixed-point representations, with minimal loss of model accuracy. Model distillation[16, 22, 37], on the other hand, is an approach that involves training a smaller, more efficient “student” model to mimic the behavior of a larger, more accurate “teacher” model by minimizing the difference between the student model’s output and the teacher model’s output. The distilled student model retains much of the teacher model’s accuracy while requiring significantly fewer resources. These model compression methods are orthogonal to offloading methods, because they achieve faster inference speed by modifying the model and sacrificing the accuracy of the result, while offloading realizes fast inference without loss of accuracy by scheduling the calculation tasks.

RPC optimization. RPC[8] is a communication protocol that enables one process to request a service from another process located on a remote computer, typically over a network. To improve RPC performance, several optimization strategies can be employed to achieve more efficient communication between remote processes and an overall enhancement in system performance: Batching[19] (aggregate multiple RPC calls into a single request), Asynchronous RPC[11] (decouple the request and response processing), and Caching[35] (Store the results of previous RPC calls). However, these optimization strategies are not effective in reducing the communication cost during model inference in robotic IoT. During the model inference process, the next operator is typically called after the previous operator has completed its execution, which renders Batching ineffective. While Asynchronous RPC and Caching enable the client to continue executing other tasks (subsequent operators) without waiting for the server’s response, they lack the ability to determine when to stop and obtain the correct computation results. Compared with the traditional RPC

optimization strategies, RRTO further reduces the communication cost by avoiding most operator’s corresponding RPC communication, which will be described with more details in Sec. III and can be considered as a specific co-design for RPC optimization strategy and transparent offloading system for model inference.

III. OVERVIEW

A. Workflow of RRTO

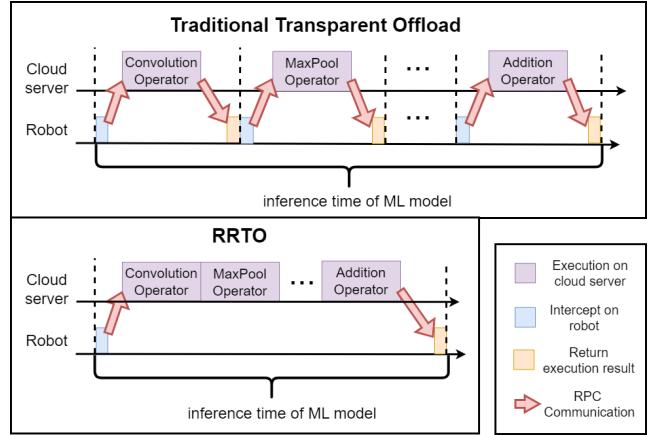


Fig. 4: Workflow of RRTO

Figure 4 illustrates the workflow of RRTO and contrasts it with traditional transparent offloading systems during model inference in robotic IoT networks. Traditional transparent offloading systems suffer from frequent RPC communication of operators, resulting in substantial communication cost and diminished system performance, including reduced GPU utilization on cloud servers, extended model inference time, and increased energy consumption per inference. The RTT communication cost for each operator is relatively minimal in data center networks, where devices connected with high-speed (e.g., 40 Gbps ~500 Gbps) networking technologies, such as InfiniBand[28] or PCIe[1]. However, in robotic IoT networks, the bandwidth between robots and powerful GPU devices is more limited and ML models commonly have hundreds of operators (e.g., 418 for [25]). For edge devices connected via Wi-Fi 6, the actual bandwidth reaches only 450 Mbps, leading to an average RTT of 2.6 milliseconds for each operator and the communication cost accounting for 98% of the total inference time in our experiments. For cloud servers, the lower bandwidth imposed by the internet further hinders performance, as evaluated in Sec. VI.

To address the communication cost issue in transparent offloading processes of ML models, RRTO introduces an automatic recording and replay mechanism. Since ML model inference can be regarded as a complex function calculation, ML models typically invoke corresponding operators (e.g., addition, convolution, max-pool) in a fixed order to obtain accurate computation results and replay these operators for each subsequent inference process. RRTO records the operators called during the first few inferences and replays the execution

of this recorded sequence, referred to as the *inference operator sequence*, for subsequent inferences.

By employing this approach, RRTO only requires the first and last operators in the inference operator sequence to be offloaded via RPC, as in traditional transparent offloading systems, to obtain the correct input and output of ML models. For the operators in the middle of the inference operator sequence, RRTO directly calls these operators on the offloading server side without requiring any RPC communication from the offloading client side, thus avoiding the inherent communication cost associated with these operators.

Notice that, although there has been substantial work on optimizing RPC communication[], RRTO goes a step further by directly eliminating the RPC communication of operators in the middle of the sequence. While existing RPC optimization methods still wait for RPCs from the offloading client to instruct the offloading server on the subsequent functions to be executed, RRTO preemptively calls the corresponding operators' functions on the offloading server side.

B. Architecture of RRTO

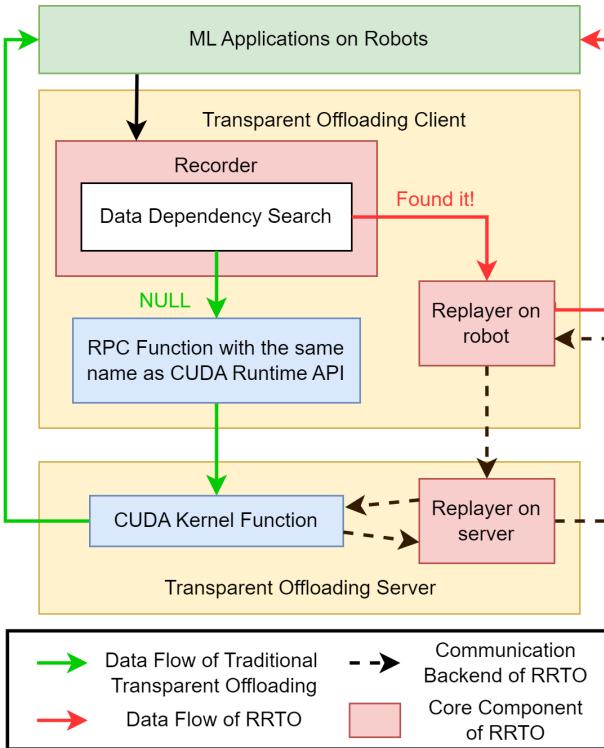


Fig. 5: Architecture of RRTO

Figure 5 presents the architecture of RRTO. In comparison with Figure 3, it is evident that RRTO implements its record/replay mechanism based on the core components of existing transparent offloading systems, and retains transparency to upper-layer applications, meaning that RRTO still does not require modifications to the source code to enable offloading.

Upon identifying and intercepting CUDA kernel function calls of operators from upper-layer ML applications, RRTO

first records the function called by the operator, including the required parameters and the return value, using its recorder. Next, RRTO attempts to find the inference operator sequence through data dependency search, which will be described with more details in Sec. IV-B. If the recorder cannot find the inference operator sequence during the first few inferences, RRTO follows the same execution pattern as traditional transparent offloading systems by offloading the execution of the operator to the cloud server via RPC, as depicted by the green lines in Figure 5.

Once the recorder identifies the inference operator sequence, RRTO initiates the replaying of the execution of the inference operator sequence for subsequent inferences using the replayer on both the robot and server, as illustrated by the red lines in Figure 5. Similar to Caching (described in Sec. II-B) in existing RPC optimization methods, the replayer on the robot returns the execution results of previous RPC calls to the upper-layer applications, allowing the offloading client to continue execution until it is blocked at the ending operator to receive the final computation result from the offloading server. Meanwhile, the offloading server replays the execution of the inference operator sequence sent by the client and returns the final computation result to the offloading client. In this manner, RRTO achieves a communication cost nearly equivalent to that of untransparent offloading methods, as it needs to transmit almost the same input and output as these methods.

Notice that, the scheduling approach utilized above offloads all computation of model inference to the cloud server. Existing computation offloading systems, such as MCOP[38], adopt more varied and flexible scheduling approaches. Depending on the model's workload, network bandwidth, and the computing power of the cloud server, these systems adaptively allocate a portion of the computation locally (compute on robots) and a portion to the cloud server. This enables faster inference by transmitting intermediate computation results with smaller transmission volumes rather than the raw input. Some works[5, 40, 26] have gone further by achieving a tradeoff between energy consumption and inference times. RRTO provides these scheduling approaches with a lower-level and finer-granularity scheduling method while maintaining nearly the same communication cost. However, this is a completely new field and is not the focus of this article. So we consider it as future work and implement the same scheduling approach in RRTO as in MCOP by adaptively choosing which operators to compute locally and which to offload to the cloud server within RRTO's recorder.

IV. DESIGN

A. Record/Replay Mechanism

Here we present how RRTO achieves record/replay mechanism. The transparent offloading client part is given in Algo. 1 and the server part is given in Algo. 2. Details of data dependency search are mainly described in the next subsection IV-B.

In Algo. 1, on the offloading client side, RRTO takes as input a CUDA kernel function called by the corresponding operator and the required parameters. The algorithm first

Algorithm 1: RRTO_on_Client

```
Input: Cuda kernel function called by the corresponding
       operator func and the required parameters args
Output: The execution result ret
Data: inference operator sequence IOS =  $\emptyset$ 
1 if IOS.empty() then
2   // recorder
3   SendRPCtoServer(func, args)
4   IOS = DataDependencySearch(func, args)
5   ret = GetRPCExecutionResult()
6   RecordReturn(ret)
7 end
8 else
9   // replayer on robot
10  if func == IOS.start()["func"] then
11    ret = StartRRTO(args, IOS)
12    // start a new inference
13  end
14  else if func == IOS.end()["func"] then
15    ret = WaitingForRRTO()
16    // Waiting for the final computation
17    result
18  end
19 end
20 else
21   | ret = IOS.find(func)["ret"]
22 end
23
24 return ret
```

checks whether the recorder has already identified the inference operator sequence (line 1). If the sequence has not been found, RRTO proceeds with the recorder phase, which includes sending an RPC to the server (line 2), performing a data dependency search (line 3), and obtaining and recording the RPC execution result (lines 4, 5). To enhance system efficiency, RRTO overlaps the data dependency search with the execution of the RPC, allowing the *DataDependencySearch* algorithm calculation to be completed while the client awaits the RPC execution result. If the inference operator sequence has been identified, RRTO proceeds with the replayer phase on the robot. This phase involves initiating RRTO for a new inference at the first operator (line 9), returning the execution results of previous RPC calls at the intermediate operators within the inference operator sequence (line 15), and waiting for the final computation result at the last operator (line 12).

In Algo. 2, the RRTO offloading server continuously awaits tasks from the client and returns the final execution results. If the client is still in the recorder phase, the RRTO offloading server processes RPC requests in the same manner as traditional transparent offloading systems (lines 2, 3). Once the client enters the replayer phase on the robot, the RRTO offloading server correspondingly transitions into the replayer phase on the server, replaying the execution of the inference operator sequence identified by the client (lines 8, 9). During

Algorithm 2: RRTO_on_Server

```
Input: client task task
Data: inference operator sequence IOS =  $\emptyset$ , the execution
      result ret
1 if task == SendRPCtoServer then
2   | func, args = GetClientInput()
3   | ret = CUDARuntimeLibrary(func, args)
4 end
5 else
6   // replayer on server
7   | args, IOS = GetClientInput()
8   foreach Op  $\in$  IOS do
9     | args = RRTOFixArgs(Op["args"], ret, args)
10    | ret = CUDARuntimeLibrary(Op["func"], args)
11  end
12 end
13 SendExecutionResultBack(ret)
```

this process, RRTO needs to adjust the parameters required by the corresponding operators, which typically consist of data or addresses of the computation results from the previous operators within the current inference.

B. Algorithm of Data Dependency Search

The performance of RRTO is heavily dependent on its ability to identify the correct inference operator sequence. If the sequence is not found accurately, even with a discrepancy of just one operator more or less, RRTO will not be able to obtain the correct inference result. Identifying the inference operator sequence is challenging, as RRTO must maintain its transparency and cannot receive any hints from upper-layer applications regarding which operators are invoked for each inference. Instead, it can only rely on log records of operators for the first few inferences, where there is no direct way to determine which inference the operator is invoked by.

The pseudocode for data dependency search is provided in Algo. 3. Due to data dependencies between operators (i.e., the calculation result of the previous operator serves as input for the next operator), RRTO first records the data dependencies between operators and constructs a relationship graph (line 2). These dependencies can be established by comparing whether parameters and calculated results between operators are the same (having the same address). Then, RRTO attempts to find the inference operator sequence based on this relationship graph.

RRTO considers operators that do not depend on any other operators as starting operators (line 3), which are candidates for the first operator in the inference operator sequence. It also considers operators that do not have any operators dependent on them as ending operators (line 4), which are candidates for the last operator in the inference operator sequence. RRTO searches for the longest covering operator sequence between starting and ending operators (line 5). Finally, RRTO verifies

Algorithm 3: DataDependencySearch

```

Input: Cuda kernel function called by the last operator
       func and the required parameters args
Output: inference operator sequence IOS
Data: Log records history =  $\emptyset$  and relationship graph
       map =  $\emptyset$ 
1 history.add(func)
2 map.update(func, args)
3 StartPoses = map.startposes()
4 EndPoses = map.endposes()
5 Sequence =
   FindLongestPair(StartPoses, EndPoses)
6 if Verify(history, Sequence) then
7   return Sequence
8 end
9 else
10  return NULL
     // cannot find
11 end

```

whether such an operator sequence can constitute a complete model inference process (line 6) by checking if the entire log records of operators can be covered by repeating this sequence.

During our implementation, we discovered that the final computation result needs to be sent back from GPU to CPU, and the ending operator can be identified by comparing the address of the final computation result. This specific observation aids RRTO in identifying the correct ending operator and dramatically reduces its search space. However, this does not entirely eliminate the need for establishing the relationship graph, because when the model has multiple outputs, an inference will involve multiple data copies from GPU to CPU and still require the Data Dependency Search algorithm to find the correct ending operator.

V. IMPLEMENTATION

We implemented RRTO within Cricket’s codebase[9], a transparent offloading system that provides a virtualization layer for CUDA applications, enabling remote execution without the need for recompiling applications. RRTO employs the same Remote Procedure Call (RPC) for communication operations as Cricket: Libtirpc[8], a transport-independent RPC library for Linux. RRTO integrates its recorder and replayers into the corresponding RPC functions in Cricket and adapts MCOP’s scheduling approach[38] by refining the scheduling granularity from MCOP’s layers of ML models to the more fine-grained operators.

VI. EVALUATION

Testbed. The evaluation was conducted on a custom four-wheeled robot (Fig 6), equipped with a Jetson Xavier NX[29] 8G onboard computer serving as the ROS master. The system runs Ubuntu 18.04 and utilizes a SanDisk 256G memory card, with ROS2 Galactic installed for application development and a dual-band USB network card (MediaTek MT76x2U) for

wireless connectivity. The Jetson Xavier NX interfaces with a Leishen N10P LiDAR, ORBBEC Astra depth camera, and an STM32F407VET6 controller via USB serial ports. Both LiDAR and depth cameras facilitate environmental perception, enabling autonomous navigation, obstacle avoidance, and SLAM mapping. The host computer processes environmental information in ROS2 Galactic, performing path planning, navigation, and obstacle avoidance before transmitting velocity and control data to corresponding ROS topics. The controller then subscribes to these topics, executing robot control tasks.

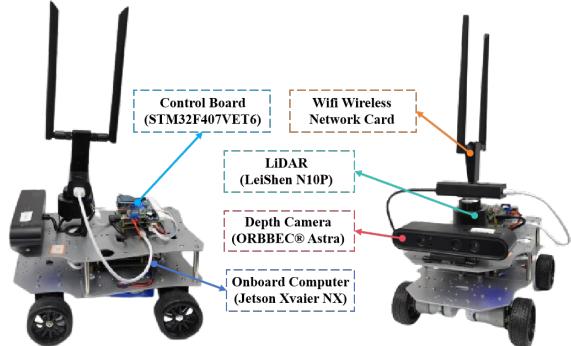


Fig. 6: The detailed composition of the robot platform

	inference	communication	standby
Power (W)	13.35	4.25	4.04

TABLE I: Power (Watt) of our robot in different states.

We documented the overall on-board energy consumption (excluding motor energy consumption for robot movement) of the robot in various states, as presented in Table I. These states include: inference, which refers to model inference with the full utilization of GPU and encompasses the energy consumption of both the CPU and GPU; communication, which involves communication with the server and includes the energy consumption of the wireless network card; and standby, during which the robot has no tasks to execute.

We evaluated two prevalent offloading scenarios for ML applications on robots, referred to as edge and cloud scenarios. In the edge scenario, computation is offloaded to a edge device, which is a PC equipped with 8xIntel(R) Core(TM) i7-7700K CPU @ 4.20GHz and NVIDIA GeForce GTX 1080 Ti 11GB, connected to our robot via Wi-Fi 6 with an average bandwidth of 450 Mbps over 160MHz channel at 5GHz frequency in our experiments. In the cloud scenario, computation is offloaded to a cloud server, which is a GPU server equipped with 48xIntel(R) Xeon(R) Silver 4116 CPU @ 2.10GHz and 4xNVIDIA GeForce RTX 2080 Ti 11GB, connected to our robot via the Internet with an average bandwidth of 160 Mbps in our experiments.

Real-world Robotic application. We evaluated a real-time people-tracking robotic application on our robot as depicted in Fig 7. To achieve seamless tracking of individuals, a minimum

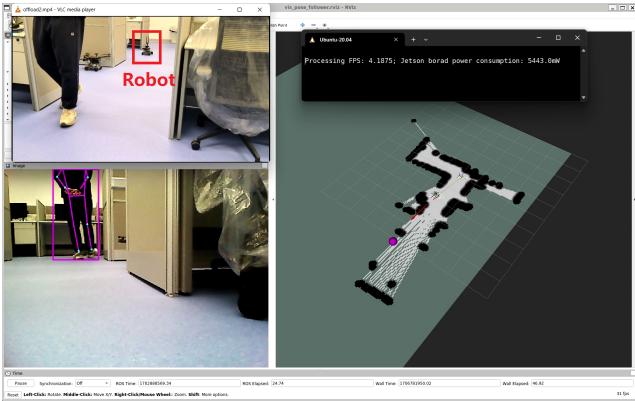


Fig. 7: A screenshot of our real-world experiment. The upper right corner displays real-time FPS and on-board energy consumption, the lower right corner shows the map created by the robot using its LiDAR, the lower left corner features the real-time view from the robot’s camera, and the upper left corner provides a third-angle observation of the entire experimental process.

detection frequency of 1 Hz for the target person is required, illustrated by the black dotted line in Fig 2a and Fig 9a. The detailed workflow is described as follows: The ORBBEC Astra depth camera on our robot generates both RGB images and corresponding depth images. First, we obtain a person’s key points in the RGB image using a well-known human pose estimation model, KAPAO[25]. Then, by utilizing the depth values corresponding to these key points in the depth image, the points are mapped to a three-dimensional map constructed by the robot’s LiDAR. A Kalman filter is applied to filter out noise and obtain a more accurate position of the person. Finally, the STM32F407VET6 controller directs the robot to the target position, enabling real-time tracking of the person.

Baselines. We compared RRTTO with local computation, MCOP[38] (a SOTA untransparent offloading system) and Cricket[9] (a SOTA transparent offloading system) when offloading computation to different GPU devices (edge devices with high bandwidth and cloud servers with limited bandwidth).

The evaluation questions are as follows:

- RQ1: How does RRTTO benefit real-world robotic applications compared to baseline systems in terms of inference time and energy consumption?
- RQ2: How does RRTTO’s record/replay mechanism work?
- RQ3: How sensitive is RRTTO to various robotic IoT environments (the bandwidth to the GPU devices and the computing power of the GPU devices)?
- RQ4: What are the limitations and potentials of RRTTO?

A. End-to-end Performance

Our evaluation results for the edge scenario, as presented in Fig 2, demonstrate that RRTTO provides performance comparable to MCOP for our robotic application. We further compared

RRTTO with baseline methods in terms of inference time and energy consumption per inference, as illustrated in Fig 8.

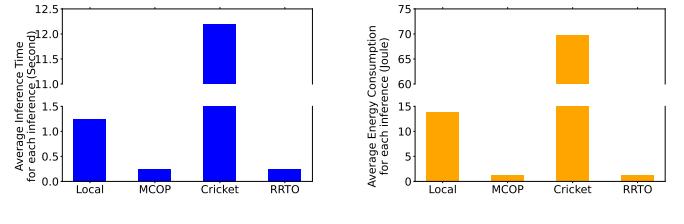


Fig. 8: Comparison between RRTTO and the baselines per inference in the edge scenario.

In terms of inference time, RRTTO achieved an 80% reduction compared to local computation and a 98% reduction compared to Cricket, as shown in Fig 8a, leading to a higher frequency in Fig 2a. Despite the additional transfers required, the powerful GPU’s shorter computation time allowed RRTTO to perform inferences faster than local computation. The significant communication cost incurred by Cricket’s transparent offloading mechanism dramatically slowed down its inference time, which will be discussed in more detail in Sec. VI-B. RRTTO successfully reduced this extremely heavy communication cost using its record/replay mechanism, achieving a similar inference time to MCOP with nearly the same communication cost.

In terms of energy consumption, RRTTO saved 90% compared to local computation and 98% compared to Cricket, as shown in Fig 8b. Although RRTTO only saved 53% power compared to local computation and 8% power compared to Cricket in Fig 2b, the shorter inference time in Fig 8a allowed RRTTO to consume significantly less energy per inference.

Notice that, the average power values in Fig 2b are not equal to those in Table I. Our application cannot fully utilize the GPU of the Jetson Xavier NX, resulting in the average power of local computation being lower than during the inference stage. Additional CPU computing tasks, such as robot control, cause the average power of offloading methods to be higher than during the communication and standby stages. Furthermore, the extremely frequent RPC function calls of Cricket generated 8% more power consumption on the CPU, and the significant communication cost caused Cricket to spend excessive time in the communication stage, wasting 98% of energy for each inference.

B. Micro-Event Analysis

To obtain a deeper understanding of the performance improvement achieved by RRTTO, we performed an analysis of the RPC function calls generated by Cricket during a single inference, which are attributable to the traditional transparent offloading mechanism. The detailed breakdown of these calls is presented in Table II.

In Table II, we observed that a significant portion, specifically 90.35%, of the RPC function calls were attributed to

CUDA Runtime API	Numbers
cudaGetDevice	3677 (80.35%)
cudaGetLastError	458 (10.00%)
cudaLaunchKernel	418 (9.13%)
cudaStreamSynchronize	6 (0.13%)
cudaMalloc	5 (0.11%)
cudaStreamIsCapturing	5 (0.11%)
cudaMemcpyHtoD	5 (0.11%)
cudaMemcpyDtoH	1 (0.02%)
cudaMemcpyDtoD	1 (0.02%)

TABLE II: Composition of RPC function calls during once inference of KAPAO

cudaGetDevice and cudaGetLastError. These calls, generated by PyTorch[32] due to our application’s reliance on this framework, serve the purpose of determining the data’s location and are essential for executing computations across multiple GPUs and parallel tasks. Even when restricting PyTorch to utilize only a single GPU sequentially and employing Caching (described in Sec. II-B) to avoid invoking these RPC functions, transparent offloading systems can only achieve inference times similar to local computation according to our experiments. However, they still cannot match the inference times achieved by MCOP or RRTTO. This is evident from the fact that cudaLaunchKernel still accounts for 9.13% of the total RPC function calls, as it informs the server about subsequent computing tasks, such as additional convolution or maxpool operations. While existing RPC optimization methods rely on waiting for RPCs of cudaLaunchKernel from the client to instruct the server regarding the subsequent computing tasks, RRTTO records these cudaLaunchKernel function calls and directly executes the subsequent computing tasks on the server without the need for communication with the client.

Regarding the remaining RPC functions, namely cudaMemcpySynchronize, cudaMalloc, cudaStreamIsCapturing, and cudaMemcpyDtoD, which collectively account for 0.37% of the total RPC calls, they primarily handle data transmission and synchronization within the GPU and can also be replayed by RRTTO on the server. However, cudaMemcpyHtoD and cudaMemcpyDtoH, accounting for 0.13% of the total RPC calls, are primarily used for data transmission between the CPU and GPU, which are mainly employed for the input and output of the ML model and cannot be replayed by RRTTO.

	MCOP	Cricket	RRTO
RPCs for each inference	N\A	4576	6
GPU utilization on edge device	29.0%	1.1%	27.5%

TABLE III: Comparison between RRTTO and the baselines on the third-party explicit features

To provide an additional perspective on the performance gain achieved by RRTTO, we compared it with the baselines using third-party explicit features, as shown in Table III. MCOP synchronizes data (input and output of the ML model

via cudaMemcpyHtoD and cudaMemcpyDtoH in Cricket) at the application layer by modifying the source code instead of using RPC. Cricket incurred a higher communication cost, leading to lower GPU utilization on the edge device. Although RRTTO still needs to handle cudaMemcpyHtoD and cudaMemcpyDtoD as Cricket does, resulting in 6 RPCs per inference, the benefits of RRTTO in terms of performance improvement are evident.

C. Sensitivity Studies

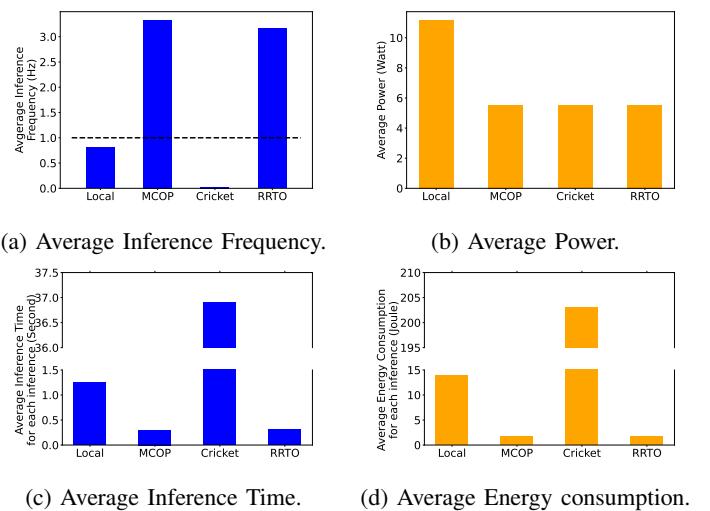


Fig. 9: Comparison between RRTTO and the baselines in the cloud scenario.

We conducted an extensive empirical evaluation of RRTTO’s performance across various robotic IoT environments, the cloud scenario illustrated in Fig 9. While cloud servers possess more powerful GPUs, their limited bandwidth causes offloading methods to incur greater communication time, ultimately resulting in slower inference times. In the cloud scenario, Cricket experiences additional performance degradation due to the significant proportion of the communication cost caused by its transparent offloading mechanism. However, with a transmission cost similar to that of MCOP, RRTTO still manages to achieve nearly same high performance as MCOP in the cloud scenario.

D. Lessons learned

Changing network bandwidth. During the implementation and evaluation of RRTTO, we discovered that the performance of computation offloading is heavily reliant on the network bandwidth of robotic IoT networks. Unlike GPU clusters equipped with fast interconnects such as InfiniBand[28], robotic IoT often lack fast and stable network connections for computation offloading. Numerous factors may cause network fluctuations, including wireless signal interference for Wi-Fi (due to obstacles[1], robot movement[12], channel preemption by other wireless devices[10], etc.) and network congestion for Internet connections (caused by multiple users accessing

simultaneously). The fine-grained offloading schedules for operators in RRTO, as described in Sec. III-B, have the potential to handle such random and violent fluctuations in network bandwidth, and we leave it as future work.

Fixed calculation logic. The record/replay mechanism of RRTO is based on the fact that operators invoked by ML models are typically fixed in order. On one hand, this mechanism allows RRTO to support other computation tasks with fixed calculation logic, even if they are not ML applications but rather computation-intensive tasks[4, 6]. On the other hand, RRTO cannot support computation tasks with unfixed calculation logic. In such tasks, the logic of each calculation differs, making it impossible for RRTO to replay them. However, computation tasks with complex logic and branching are better suited for execution in CPU rather than GPU[33], let alone RRTO.

Future work. We would like to apply and evaluate RRTO in a wider variety of real-world applications on robots in the future. Also, it is of interest to explore further improvements of RRTO such as building an edge computing power network that offloads computation to other idle robots or edge devices[2] through fine-grained job scheduling at the operator level to ensure full utilization of GPU resources. We believe that such investigation will enable even faster, more energy-efficient, and more robust inference capabilities for ML applications in real-world robotic IoT networks.

VII. CONCLUSION

In this paper, we introduce RRTO, a high-performance transparent offloading system designed for model inference on robotic IoT. RRTO tackles the inherent communication cost associated with traditional transparent offloading mechanisms by introducing a novel record/replay mechanism, achieving the same high performance as the SOTA untransparent offloading method without modifying any source code. We envision that RRTO will significantly streamline the deployment of a diverse array of ML applications on mobile robots in real-world settings. By offering fast and energy-efficient inference capabilities, RRTO enables these robots to execute complex tasks with heightened efficiency and effectiveness.

ACKNOWLEDGMENTS

REFERENCES

- [1] Badrul Alam, Giovanna Calò, Gaetano Bellanca, Jacopo Nanni, Ali Emre Kaplan, Marina Barbiroli, Franco Fuschini, Paolo Bassi, Jinous Shafiei Dehkordi, Velio Tralli, et al. Numerical and experimental analysis of on-chip optical wireless links in presence of obstacles. *IEEE Photonics Journal*, 13(1):1–11, 2020.
- [2] Yixin Bao, Yanghua Peng, Chuan Wu, and Zongpeng Li. Online job scheduling in distributed machine learning clusters. In *IEEE INFOCOM 2018-IEEE Conference on Computer Communications*, pages 495–503. IEEE, 2018.
- [3] Anh-Quan Cao and Raoul de Charette. Monoscene: Monocular 3d semantic scene completion. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 3991–4001, 2022.
- [4] Niels Cautaerts, Phillip Crout, Håkon W Ånes, Eric Prestat, Jiwon Jeong, Gerhard Dehm, and Christian H Liebscher. Free, flexible and fast: Orientation mapping using the multi-core and gpu-accelerated template matching capabilities in the python-based open source 4d-stem analysis toolbox pyxem. *Ultramicroscopy*, 237:113517, 2022.
- [5] Xing Chen, Jianshan Zhang, Bing Lin, Zheyi Chen, Katinka Wolter, and Geyong Min. Energy-efficient offloading for dnn-based smart iot systems in cloud-edge environments. *IEEE Transactions on Parallel and Distributed Systems*, 33(3):683–697, 2021.
- [6] Rohit Chowdhury and Deepak Subramani. Optimal path planning of autonomous marine vehicles in stochastic dynamic ocean flows using a gpu-accelerated algorithm. *IEEE Journal of Oceanic Engineering*, 47(4):864–879, 2022.
- [7] Alexandre Défossez, Yossi Adi, and Gabriel Synnaeve. Differentiable model compression via pseudo quantization noise. *arXiv preprint arXiv:2104.09987*, 2021.
- [8] Steve Dickson. libtirpc: Transport independent rpc library. <https://git.linux-nfs.org/?p=steved/libtirpc.git>, 2024. Git repository.
- [9] Niklas Eiling, Jonas Baude, Stefan Lankes, and Antonello Monti. Cricket: A virtualization layer for distributed execution of cuda applications with checkpoint/restart support. *Concurrency and Computation: Practice and Experience*, 34(14), 2022. doi: 10.1002/cpe.6474.
- [10] Khaled Eshteiwi, Georges Kaddoum, Bassant Selim, and François Gagnon. Impact of co-channel interference and vehicles as obstacles on full-duplex v2v cooperative wireless network. *IEEE Transactions on Vehicular Technology*, 69(7):7503–7517, 2020.
- [11] Stijn Eyerman and Ibrahim Hur. Efficient asynchronous rpc calls for microservices: Deathstarbench study. *arXiv preprint arXiv:2209.13265*, 2022.
- [12] Yange Feng, Enrico Benassi, Liqiang Zhang, Xiaojuan Li, Daoai Wang, Feng Zhou, and Weimin Liu. Concealed wireless warning sensor based on triboelectrification and human-plant interactive induction. *Research*, 2021.
- [13] Stefan Gheorghe and Mihai Ivanovici. Model-based weight quantization for convolutional neural network compression. In *2021 16th International Conference on Engineering of Modern Electric Systems (EMES)*, pages 1–4. IEEE, 2021.
- [14] Giulio Giunta, Raffaele Montella, Giuseppe Agrillo, and Giuseppe Coviello. A gpgpu transparent virtualization component for high performance computing clouds. In *Euro-Par 2010-Parallel Processing: 16th International Euro-Par Conference, Ischia, Italy, August 31-September 3, 2010, Proceedings, Part I 16*, pages 379–391. Springer,

- 2010.
- [15] Cheng Gong, Yao Chen, Ye Lu, Tao Li, Cong Hao, and Deming Chen. Vecq: Minimal loss dnn model compression with vectorized weight quantization. *IEEE Transactions on Computers*, 70(5):696–710, 2020.
 - [16] Jianping Gou, Baosheng Yu, Stephen J Maybank, and Dacheng Tao. Knowledge distillation: A survey. *International Journal of Computer Vision*, 129:1789–1819, 2021.
 - [17] Vinayak Honkote, Dileep Kurian, Sriram Muthukumar, Dibyendu Ghosh, Satish Yada, Kartik Jain, Bradley Jackson, Ilya Klotchkov, Mallikarjuna Rao Nimmagadda, Shreela Dattawadkar, et al. 2.4 a distributed autonomous and collaborative multi-robot system featuring a low-power robot soc in 22nm cmos for integrated battery-powered minibots. In *2019 IEEE International Solid-State Circuits Conference-(ISSCC)*, pages 48–50. IEEE, 2019.
 - [18] K J Joseph, Salman Khan, Fahad Shahbaz Khan, and Vineeth N Balasubramanian. Towards open world object detection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 5830–5840, June 2021.
 - [19] Nikita Lazarev, Shaojie Xiang, Neil Adit, Zhiru Zhang, and Christina Delimitrou. Dagger: efficient and fast rpcs in cloud microservices with near-memory reconfigurable nics. In *Proceedings of the 26th ACM International Conference on Architectural Support for Programming Languages and Operating Systems*, pages 36–51, 2021.
 - [20] Qingbiao Li, Fernando Gama, Alejandro Ribeiro, and Amanda Prorok. Graph neural networks for decentralized multi-robot path planning. In *2020 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 11785–11792. IEEE, 2020.
 - [21] Yiming Li, Zhiding Yu, Christopher Choy, Chaowei Xiao, Jose M Alvarez, Sanja Fidler, Chen Feng, and Anima Anandkumar. Voxformer: Sparse voxel transformer for camera-based 3d semantic scene completion. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 9087–9098, 2023.
 - [22] Tao Lin, Lingjing Kong, Sebastian U Stich, and Martin Jaggi. Ensemble distillation for robust model fusion in federated learning. *Advances in Neural Information Processing Systems*, 33:2351–2363, 2020.
 - [23] Yu-Shiang Lin, Chun-Yuan Lin, Che-Rung Lee, and Yeh-Ching Chung. qcuda: Gpgpu virtualization for high bandwidth efficiency. In *2019 IEEE International Conference on Cloud Computing Technology and Science (CloudCom)*, pages 95–102. IEEE, 2019.
 - [24] Shuai Liu, Xin Li, Huchuan Lu, and You He. Multi-object tracking meets moving uav. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 8876–8885, June 2022.
 - [25] William McNally, Kanav Vats, Alexander Wong, and John McPhee. Rethinking keypoint representations: Modeling keypoints and poses as objects for multi-person human pose estimation. *arXiv preprint arXiv:2111.08557*, 2021.
 - [26] Thaha Mohammed, Carlee Joe-Wong, Rohit Babbar, and Mario Di Francesco. Distributed inference acceleration with adaptive dnn partitioning and offloading. In *IEEE INFOCOM 2020-IEEE Conference on Computer Communications*, pages 854–863. IEEE, 2020.
 - [27] NVIDIA. cuda runtime apis. <https://docs.nvidia.com/cuda/cuda-runtime-api/index.html>, 2024.
 - [28] NVIDIA. Infiniband networking solutions. <https://www.nvidia.com/en-us/networking/products/infiniband/>, 2024.
 - [29] NVIDIA. The world’s smallest ai supercomputer. <https://www.nvidia.com/en-us/autonomous-machines/embedded-systems/jetson-xavier-series/>, 2024.
 - [30] Takeshi Ohkawa, Kazushi Yamashina, Hitomi Kimura, Kanemitsu Ootsu, and Takashi Yokota. Fpga components for integrating fpgas into robot systems. *IEICE TRANSACTIONS on Information and Systems*, 101(2):363–375, 2018.
 - [31] Brian Plancher, Sabrina M. Neuman, Thomas Bourgeat, Scott Kuindersma, Srinivas Devadas, and Vijay Janapa Reddi. Accelerating robot dynamics gradients on a cpu, gpu, and fpga. *IEEE Robotics and Automation Letters*, 6(2):2335–2342, 2021. doi: 10.1109/LRA.2021.3057845.
 - [32] pytorch. pytorch. <https://www.pytorch.org>, 2024.
 - [33] Viktor Rosenfeld, Sebastian Breß, and Volker Markl. Query processing on heterogeneous cpu/gpu systems. *ACM Computing Surveys (CSUR)*, 55(1):1–38, 2022.
 - [34] Matthias Schneider, Fabian Haag, Abdul Kader Khalil, and David Albert Breunig. Evaluation of communication technologies for distributed industrial control systems: Concept and evaluation of 5g and wifi 6. *Procedia CIRP*, 107:588–593, 2022. ISSN 2212-8271. doi: <https://doi.org/10.1016/j.procir.2022.05.030>. URL <https://www.sciencedirect.com/science/article/pii/S2212827122003146>. Leading manufacturing systems transformation – Proceedings of the 55th CIRP Conference on Manufacturing Systems 2022.
 - [35] Arjun Singhvi, Aditya Akella, Maggie Anderson, Rob Cauble, Harshad Deshmukh, Dan Gibson, Milo MK Martin, Amanda Strominger, Thomas F Wenisch, and Amin Vahdat. Cliquemap: Productionizing an rma-based distributed caching system. In *Proceedings of the 2021 ACM SIGCOMM 2021 Conference*, pages 93–105, 2021.
 - [36] Binyu Wang, Zhe Liu, Qingbiao Li, and Amanda Prorok. Mobile robot path planning in dynamic environments through globally guided reinforcement learning. *IEEE Robotics and Automation Letters*, 5(4):6932–6939, 2020. doi: 10.1109/LRA.2020.3026638.
 - [37] Lin Wang and Kuk-Jin Yoon. Knowledge distillation and student-teacher learning for visual intelligence: A review and new outlooks. *IEEE transactions on pattern analysis and machine intelligence*, 44(6):3048–3068, 2021.
 - [38] Huaming Wu, William J. Knottenbelt, and Katinka Wolter. An efficient application partitioning algorithm in mobile environments. *IEEE Transactions on Parallel*

- and Distributed Systems*, 30(7):1464–1480, 2019. doi: 10.1109/TPDS.2019.2891695.
- [39] Zhaoyang Xia, Youquan Liu, Xin Li, Xinge Zhu, Yuexin Ma, Yikang Li, Yuenan Hou, and Yu Qiao. Scpnet: Semantic scene completion on point cloud. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 17642–17651, 2023.
 - [40] Min Xue, Huaming Wu, Guang Peng, and Katinka Wolter. Ddpqn: An efficient dnn offloading strategy in local-edge-cloud collaborative environments. *IEEE Transactions on Services Computing*, 15(2):640–655, 2021.
 - [41] Yang Yang, Li Juntao, and Peng Lingling. Multi-robot path planning based on a deep reinforcement learning dqn algorithm. *CAAI Transactions on Intelligence Technology*, 5(3):177–183, 2020.